

Alek Westover
Old Brainstorming

1 chill

Algorithm 1 Alg greed

```
1: procedure GREED0
2:   while True do
3:     if all processors are idle then
4:       if If there are more than  $p/k$  queued tasks then
5:         schedule all (or  $p$ ) in serial
6:       else
7:         Schedule one in parallel
8:       end if
9:     end if
10:  end while
11: end procedure
```

Algorithm 2 Alg greed

```
1: procedure GREED1
2:   while True do
3:     if all processors are idle then
4:       if If there are more than  $p/k$  queued tasks then
5:         schedule all (or  $p$ ) in serial
6:       else
7:         Schedule all in parallel
8:       end if
9:     end if
10:  end while
11: end procedure
```

Algorithm 3 Alg greed

```
1: procedure GREED2
2:   while True do
3:     if all processors are idle then
4:        $q \leftarrow$  number of queued tasks
5:       schedule  $\lfloor q/p \rfloor$  tasks in serial to each processor
6:        $q' \leftarrow q \bmod p$ 
7:       if  $q' > p/k$  then
8:         Schedule  $q'$  tasks in serial
9:       else
10:        Schedule  $q'$  tasks in parallel (distributing work equally)
11:       end if
12:     end if
13:   end while
14: end procedure
```

You can probably do a similar thing for the non-symmetric case. The main idea is locally be greedy, but also lazy about scheduling. So in the non-symmetric case you'd probably be like "if all processors are idle, schedule in a good way."

Algorithm 4 Alg chill

```

1: procedure CHILL
2:   while True do
3:     if all processors are idle then
4:       if If there are more than  $p/k$  queued tasks then
5:         schedule all (or  $p$ ) in serial
6:       else
7:         Schedule one in parallel
8:       end if
9:     end if
10:  end while
11: end procedure

```

2 Alg X

Algorithm 5 Alg X

```

1: This procedure is continuously running.
2: procedure X
3:   if there are more than  $p/k$  queued tasks mod  $p$  then
4:     for each task running in parallel do
5:       if it has more than 1 total work left and the number of queued tasks mod  $p$  is at most  $p-1$  then
6:         kill this task it (i.e. put it on the queue)
7:       end if
8:     end for
9:     Schedule as many queued tasks as possible to processors that have less than 1 work assigned to them.
    (scheduling to minimize backlog, i.e. scheduling in ascending order of backlog)
10:  end if
11:  if There is an idle processor then
12:    if There is a queued task then
13:      if backlog  $\geq 1$  then
14:        schedule tasks in serial on any idle processors
15:      end if
16:      if backlog  $< 1$  then
17:        schedule a task in parallel, scheduling as balancedly as possible
18:      end if
19:    end if
20:    if There is no queued task and there is a serial task that could be cancelled and then redistributed
    that would result in all cups that are getting the redistribution stuff end up with less work than the thing
    that was cancelled from then
21:      Cancel the serial task and reschedule as specified
22:    end if
23:  end if
24: end procedure

```

Claim 1. *Alg X is good.*

Proof. This seems basically impossible to prove, it's a super complicated algorithm with so much branching. \square

3 some more strategies

- randomized (smoothed analysis)

- look at discrete version?

Algorithm 6 Randomized alg

```
1: This procedure is continuously running.
2: procedure RANDOMIZEDSTRATEGY
3:
4:   if backlog would be made smaller by cancelling everything and swapping the mode, and scheduling
     everything according to the new mode then
5:     do it!
6:   end if
7: end procedure
```

Algorithm 7 Alg binary

```
1: This procedure is continuously running.
2: procedure BINARY(mode)
3:   when you get a new task schedule it to minimize backlog in the current mode (mode is serial or parallel)
4:   if backlog would be made smaller by cancelling everything and swapping the mode, and scheduling
     everything according to the new mode then
5:     do it!
6:   end if
7: end procedure
```
